CUDA Notes

Eric Andrews

March 8, 2018

CUDA Implementation

- General Structure

    - Send matrices to be multiplied to the GPU
    - Have each thread compute the element of C corresponding to its id by iterating through the corresponding row and column